

03 - Models

- Models
 - Physical
 - Architectural
 - Fundamental
- Coulouris 1
- Coulouris 2
- Saltzer_84.pdf

Physical models

- Early Distributed Systems
 - Late 70/early 80s
 - 10..100 nodes
 - Local network / Homogeneous / Few services
- Internet scale
 - 90s
 - Internet based / Network of networks / Static
 - Global / Heterogeneous (but server or client) / Open
- Contemporary
 - Mobile nodes (Wifi, GSM)
 - Ubiquitous (embedded in objects and environment)
 - Systems of systems (cloud)
 - New level of heterogeneity (architecture and capabilities)

Physical models

<i>Distributed systems:</i>	<i>Early</i>	<i>Internet-scale</i>	<i>Contemporary</i>
<i>Scale</i>	Small	Large	Ultra-large
<i>Heterogeneity</i>	Limited (typically relatively homogenous configurations)	Significant in terms of platforms, languages and middleware	Added dimensions introduced including radically different styles of architecture
<i>Openness</i>	Not a priority	Significant priority with range of standards introduced	Major research challenge with existing standards not yet able to embrace complex systems
<i>Quality of service</i>	In its infancy	Significant priority with range of services introduced	Major research challenge with existing services not yet able to embrace complex systems

Architectural models

- What entities are communicating in a DS?
 - Communication entities
- How they communicate?
 - Communication paradigms
- What roles and responsibilities they have?
- How are they mapped on the physical infrastructure?

Architectural models

<i>Communicating entities (what is communicating)</i>		<i>Communication paradigms (how they communicate)</i>		
<i>System-oriented entities</i>	<i>Problem- oriented entities</i>	<i>Interprocess communication</i>	<i>Remote invocation</i>	<i>Indirect communication</i>
Nodes	Objects	Message passing	Request- reply	Group communication
Processes	Components	Sockets	RPC	Publish-subscribe
	Web services	Multicast	RMI	Message queues
				Tuple spaces
				DSM

Architectural models

Communication entities

- DS = Processes + IPC
- Processes
 - Some systems don't have processes
 - Most systems have threads
 - the real communication endpoint
- Objects
 - Migrations of OO to DS
 - OO Designs and OO programming
 - Natural decomposition unit for a problem
 - Accessed by and interface

Architectural models

Communication entities

- Components
 - Provide interface like objects
 - Specify assumption
 - Dependencies are explicit and used to pair components
 - Contracts
- Web services
 - Implementation of Objects and Components
 - On the WEB
 - Identified by a URL
 - Defined/described/discovered by XML

Architectural models

Communication paradigms

- Remote invocation
 - Request-reply protocols
 - Low level
 - Programmer creates/sends messages
 - Example: HTTP

Architectural models

Communication paradigms

- Remote invocation
 - Remote procedure calls
 - Attributed to Birrel and Nelson (84)
 - RPC system hides
 - Distribution
 - Encoding/decoding messages
 - Passing of message
 - Semantic of the procedure call
 - Supports client/server
 - Server offer set of operations (by interfaces)
 - Clients call those operations directly as if local
 - Access and location transparency
 - (minimal)

Architectural models

Communication paradigms

- Remote invocation
 - Remote method invocation
 - Resembles RPC
 - But in a world of objects
 - Client objects invoke methods on remote objects
 - Underling details are hidden
 - May support
 - Object identity
 - Pass objects as parameter
 - Tight integration to OO languages

Architectural models

Communication paradigms

- Remote invocation
 - Two-way relationship (sender \leftrightarrow receiver)
 - Receiver identity is known
 - Both parties exist simultaneously at the same time
 - Direct communication
- Indirect communication
 - Space uncoupling
 - Sender does not know who is sending to
 - Time decoupling
 - Senders and receiver do not need to exist at the same time

Architectural models

Communication paradigms

- Indirect communication
 - Group communication
 - Delivery of messages to groups of recipients
 - One-to-many communication
 - Communication relies on groups abstraction
 - Recipients join groups
 - Publish-subscribe
 - Information dissemination/ distributed events
 - Producers(publishers) distribute
 - Information items of interest (events)
 - Consumers register the interest or events to receive

Architectural models

Communication paradigms

- Indirect communication
 - Message Queues
 - Point-to-point communication channel
 - Producer places message on Queue
 - Consumer retrieves message
 - Consumers are notified of message availability
 - Tuples space
 - Communication of performed by the access to shared structured data (tuples)
 - Add tuples to the persistent tuple space
 - Consumers can read or deleted existing tuples
 - Can be client server or P2P

Architectural models

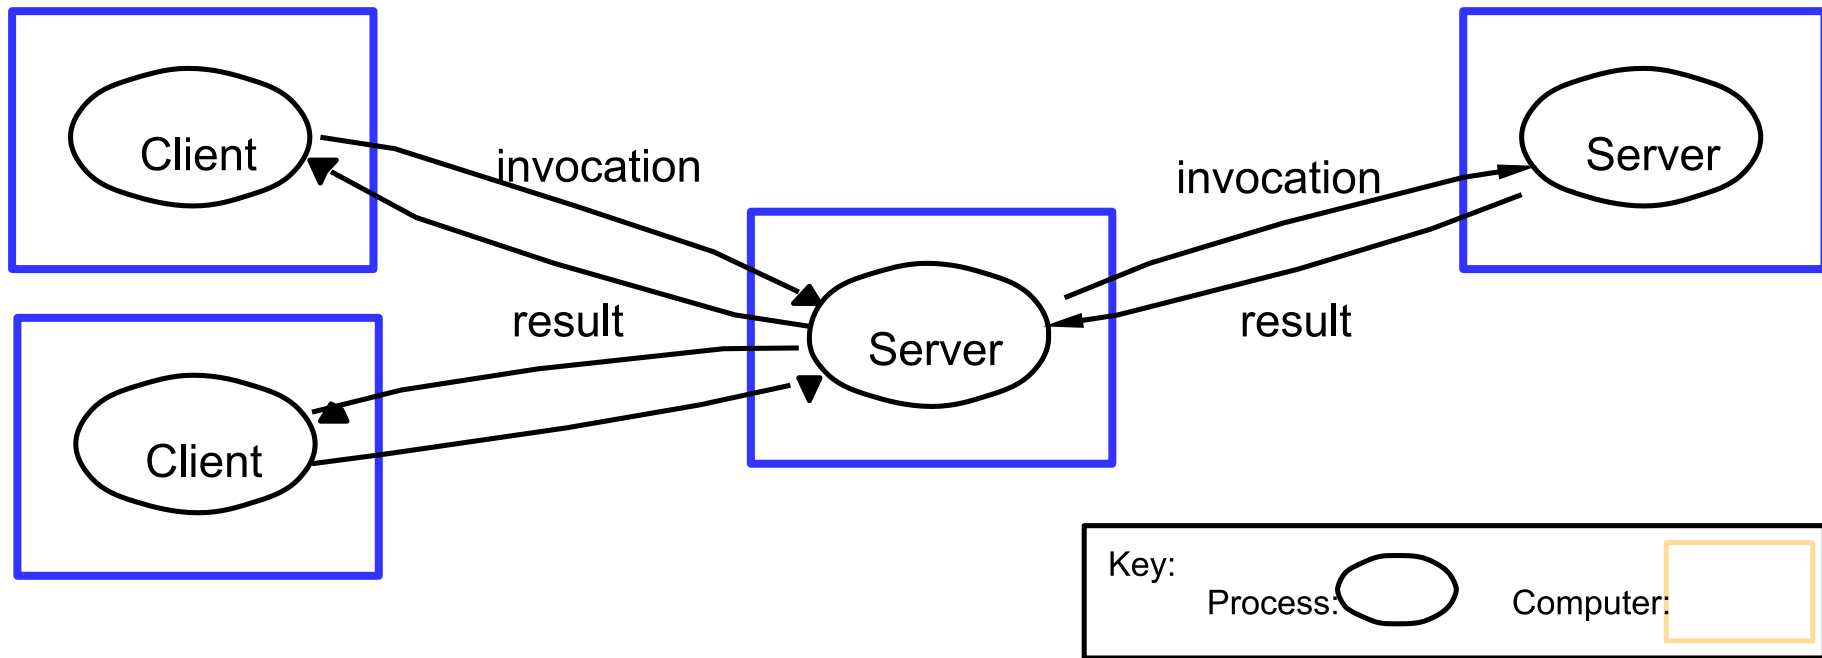
Communication paradigms

- Indirect communication
 - Distributed Shared Memory
 - DSM systems provide a view of a shared memory space
 - Composed of data on multiple remote nodes
 - Programmers are presented with the abstraction of reading/writing local memory
 - All accesses are to local address space
 - Although data can be on a remote node
 - Infrastructure guarantees
 - Copies of data are provided in a timely manner
 - Synchronization and consistency of data

Architectural models

Roles and responsibilities

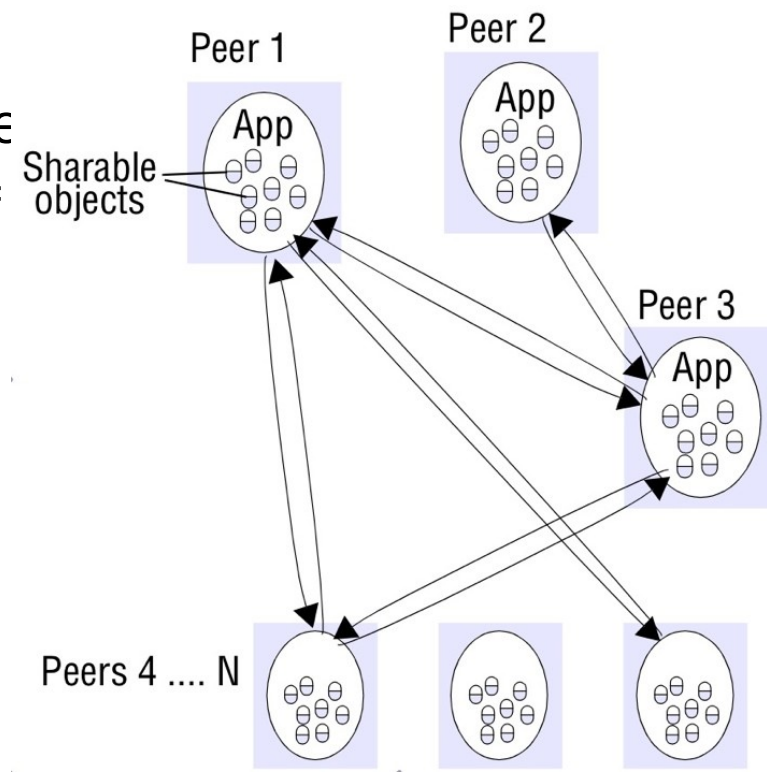
- Client-Server
 - Client processes interact with individual servers
 - (potentially) In separate hosts
 - To access shared resources
 - Servers can also be clients



Architectural models

Roles and responsibilities



- Peer-to-peer
 - All processes have similar roles
 - Interacting cooperatively as *peers*
 - No distinctions between client or server
 - Resources owned by users
 - Can be put to use to support the service
 - Resources increase with the number of users
 - Data objects are shared and distributed
 - Distribution and replication increases complexity



Architectural models

Placement

- Where to place entities on the physical model?
 - Distribute service among several servers
 - Data may be partitioned
 - Data may be replicated
 - Caching
 - Store of recently used data
 - Locally to the client
 - On a separate server
 - Objects are retrieved from cache
 - If available



Architectural models

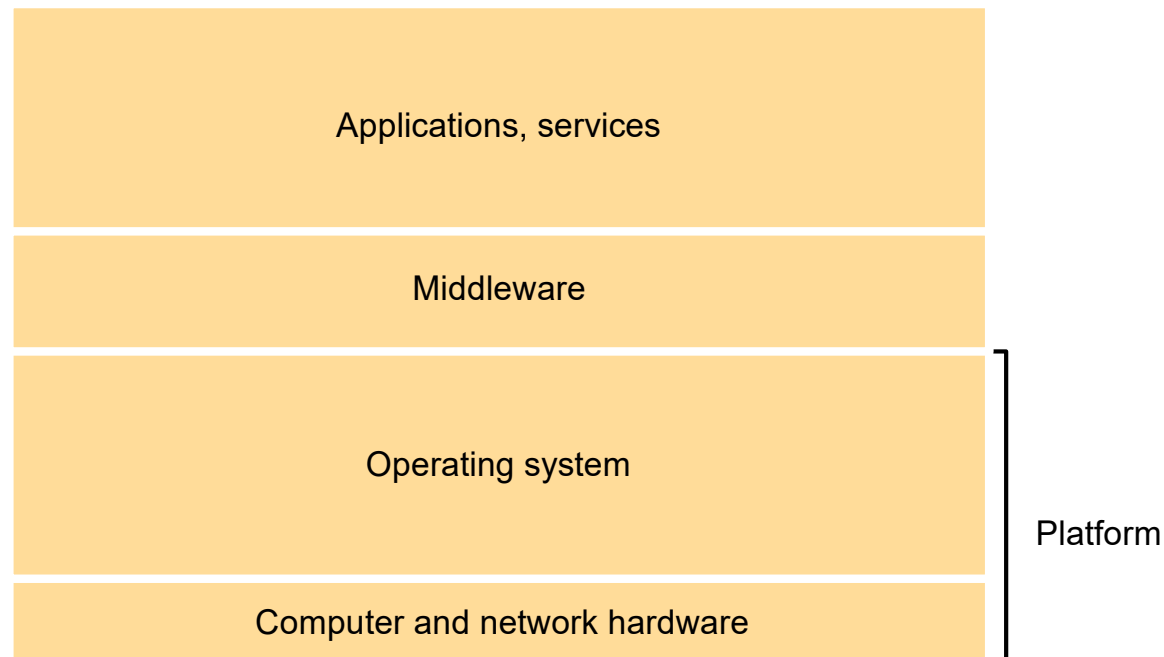
Placement

- Mobile code
 - Code is downloaded from the server
 - Executed on the client
 - No network delays
- Mobile agents
 - Code (and data) roams on computers
 - Executes on behalf of other
 - Invokes local services
 - Lower execution time
 - Code transfer + local invocation vs remote invocation

Architectural models

Architectural patterns

- Layering
 - Related to abstraction
 - Each layer offers an abstraction
 - Higher layers not aware of “lower” implementations



Architectural models

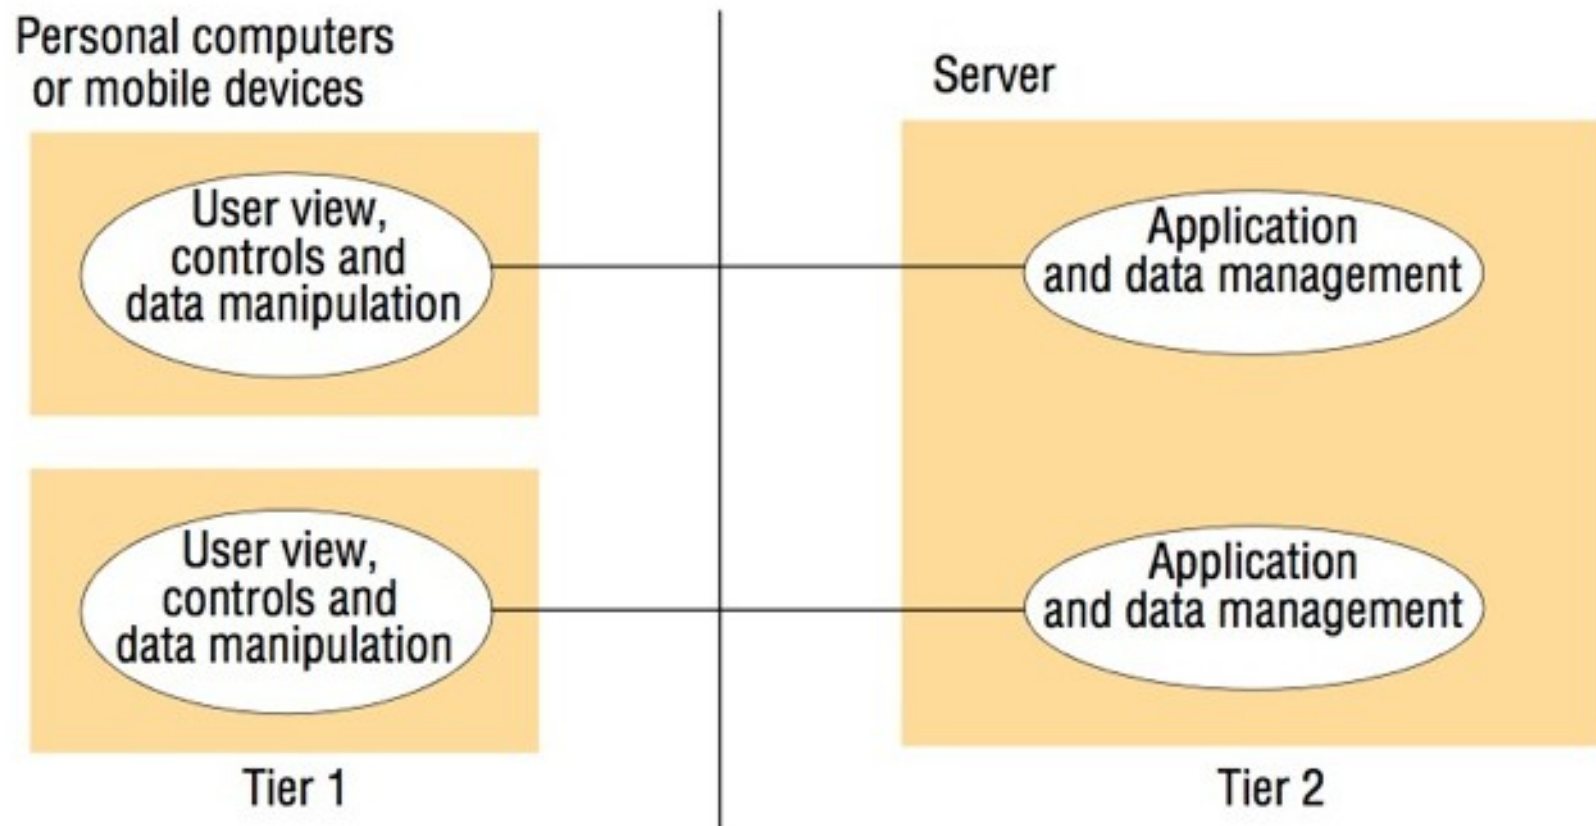
Architectural patterns

- Tiered
 - Related to composition
 - Complements layerings
 - Layering: vertical organization of services into layers
 - Tiering: Distribution of a given layer into appropriate servers
 - Presentation logic
 - Application logic
 - Data logic

Architectural models

Architectural patterns


- 2-Tiered`



Architectural models

Architectural patterns


- 3-Tiered



Architectural models

Architectural patterns

- Thin-Client
 - Moves complexity away from end-user
 - Client has no logic only presentation



Architectural models


Architectural patterns

- Proxy
 - Offers the same interface as the server
 - Located on the local client
 - Is contacted and redirect calls to remote node
 - Offers location transparency
 - Can encapsulate other functionalities
 - Placement policies of replication
 - Caching
 - authentication, ...

Architectural models

Architectural patterns

- Brokerage
 - Supporting interoperability



Fundamental models

- Describe the general and fundamental characteristics of a DS
 - Not how it is implemented
- Define the assumptions about the system
- Allows the generalization about (imp)possibilities
 - General purpose algorithms
 - Desirable properties that are guaranteed
- Interaction
- Failure
- Security

Fundamental models

Interaction

- Assumptions about the communication channels
 - Latency: delay between start of message transmission and receiving
 - Time taken on transmission
 - Delay assessing network
 - Time of message processing (on the OS)
 - Bandwidth: Amount of information that can be transmitted on a give time
 - Jitter: variation in time take to deliver a series of messages

Fundamental models

Interaction

- Assumptions about computer clocks
 - Each computer has local clock
 - Used to obtain current time
 - Simultaneous clock read render different values
 - Clocks on computers drift at different rates
 -

Fundamental models

Interaction

- Synchronous Distributed Systems
 - Time to execute each step
 - Has lower bound
 - Has higher bound
 - Messages
 - Are received within a known bounded time
 - Process clocks
 - Drift rate has known bound
 - Can be built if
 - It is possible to guarantee previous bounds

Fundamental models

Interaction

- Asynchronous Distributed Systems
 - There are no bounds for
 - Process execution speed
 - Message transmission delays
 - Clock drift rates
 - Internet
 - Good example
 - No limits to server or network load
 - Email
 - May take days to arrive
 - Servers can have drifted clock

Fundamental models

Failure

- In a DS processes and communication can fail
- Failure model defines how failures can occur
 - In order to provide understanding of their effects
- Omission failures
 - Process omission failures:
 - process does not execute the task (crashes)
 - Can be detected using timeouts
 - Communication omission failures
 - A sent message is not delivered to the receptor
 - Both benign

Fundamental models

Failure

<i>Class of failure</i>	<i>Affects</i>	<i>Description</i>
<i>Fail-stop</i>	Process	Process halts and remains halted. Other processes may detect this state.
<i>Crash</i>	Process	Process halts and remains halted. Other processes may not be able to detect this state.
<i>Omission</i>	Channel	A message inserted in an outgoing message buffer never arrives at the other end's incoming message buffer.
<i>Send-omission</i>	Process	A process completes a <i>send</i> , <i>but the message is not put in its outgoing message buffer</i> .
<i>Receive-omission</i>	Process	A message is put in a process's incoming message buffer, but that process does not receive it.
<i>Arbitrary (Byzantine)</i>	Process or Channel	Process/channel exhibits arbitrary behavior: it may send/transmit arbitrary messages at arbitrary times, commit omissions; a process may stop or take an incorrect step.

Fundamental models

Failure

- Timing failures
 - Applicable to synchronous systems
 - With limits to execution, delivery times and clock drifts

<i>Class of failure</i>	Affects	Description
Clock	Process	Process's local clock exceeds the bounds on its
Performance	Process	Process exceeds the bounds on the interval between two steps.
Performance	Channel	A message's transmission takes longer than the stated bound.

Fundamental models

Failure

- Arbitrary (Byzantine) failures
 - Any error can occur
 - A process set wrong values in data items
 - A process answers a wrong values
 - A process arbitraly fails
 - No way to distinguish Process fail from no answer :(
 - A message gets corrupted
 - A bogus messages is created
 - A messages is intercepted from delivery
 - A messages is delivered twice

Fundamental models

Masking failures

- A reliable system can be composed of
 - Unreliable components that exhibit failures
- Knowledge of each component failures
 - Allows a service to mask them
- Hiding a failure
 - e.g. retrying
- Converting to a more acceptable failure
 - Recovering an old version of the file

Fundamental models

Reliability on one-to-one communication

- Basic communication channels
 - can exhibit failures (e.g. omission)
 - But can be used to build a service that masks some failures
 - Provides reliable communication
- Reliable communication
 - Validity: any message is eventually delivered
 - Integrity: the received message is identical to the sent one

Network Platform

- Networking issues
 - Performance
 - Scalability
 - Reliability
 - Security
 - Mobility
 - QoS
- Types of networks
 - Personal Area
 - (wireless) Local Area
 - (wireless) Wide Area
 - (wireless) Metropolitan Area
- Handled by the platform
 - OS + HW
 - Low level protocols

Middleware

- SW layer that
 - Provides a programming abstraction
 - Masks underlying heterogeneity
 - Network, HW, OS, programming languages
 - Provides a uniform computational model
 - To be used by programmers of distributed applications
 - RMI, remote event notification, distributed transactions, ...

Middleware

<i>Major categories:</i>	<i>Subcategory</i>	<i>Example systems</i>
<i>Distributed objects (Chapters 5, 8)</i>	Standard	RM-ODP
	Platform	CORBA
	Platform	Java RMI
<i>Distributed components (Chapter 8)</i>	Lightweight components	Fractal
	Lightweight components	OpenCOM
	Application servers	SUN EJB
	Application servers	CORBA Component Model
	Application servers	JBoss
<i>Publish-subscribe systems (Chapter 6)</i>	-	CORBA Event Service
	-	Scribe
	-	JMS
<i>Message queues (Chapter 6)</i>	-	Websphere MQ
	-	JMS
<i>Web services (Chapter 9)</i>	Web services	Apache Axis
	Grid services	The Globus Toolkit
<i>Peer-to-peer (Chapter 10)</i>	Routing overlays	Pastry
	Routing overlays	Tapestry
	Application-specific	Squirrel
	Application-specific	OceanStore
	Application-specific	Ivy
	Application-specific	Gnutella

Middleware

Limitations

- Some application rely only on middleware
 - If suitable for client-servers can use RPC
 - Name/Address database
- Not all issues can be handled by middleware
 - Some dependability aspects
- Large file transfer over unreliable link
 - TCP offers some error detection and correction
 - TCP does not recover from major network interruption
 - If service offers a new level of fault tolerance
 - Must maintain a progress level
 - Must resumes transmission with a new TCP connection

Middleware

Limitations

- End-to-end argument
 - Saltzer et al (84)
- Some communication functions can only be completely and reliably implemented
 - With knowledge and help of the application standing at the end points of the communication system
- In the previous case:
 - TCP does not know how to restart a file transmission
 - Client must know where to restart
 - Server must receive information about restart