

06 – Naming

- What's in a name
- Naming System
- Design Tradeoffs
- Coulouris, Ch 9
- Craft, 1923
- Ahmed, 2005
- Subharthi, 2009

Naming

- Uses/affects
 - Identification
 - Location
 - Sharing
- Concerns
 - The name contains information?
 - Are names and references decoupled?
 - Can the system evolve?
 - Length
 - Syntax

Naming relevance

- What (and how) you can name
 - Fundamental to what systems do
 - Fundamental to how systems work
 - Fundamental to how systems evolve
- Uniformity
 - Naming and operations
 - Allow the operation of unknown things

Phone numbers

- Before 1999
 - Dynamic length
 - Local numbers
 - Prefix (outside calls)
 - xxx xxxx (within Lisbon)
 - 01 xxx xxxx (from the rest of Portugal)
 - +351 1 xxx xxxx (outside Portugal)
 - 0 1 YYYYYYYYYY (US)
 - 0 936 xxx xxx x

Phone numbers

- After 1999
 - fixed length (9 digits)
 - global numbers
 - 21 xxx xxxx (within Lisbon)
 - 21 xxx xxxx (from the rest of Portugal)
 - +351 21 xxx xxxx (outside Portugal)
 - 96 xxx xxxx

Phone numbers

- Differences
 - 310347
 - 21 3100347
 - 0936 990112
 - 96 9900112
- What's in the structure?
 - Local/global
 - Open/close
 - Reflects system?
- What can be accessed?
- Who generates names?
- What information is in the name?
- Are names locators?
 - Addresses?

- The nature of phone numbers determines
 - who can be called, and how easily
- Phone numbers embedded lots of information
 - Country where phone is hosted
 - Well known numbers have special use (911)
- Assignment is delegated to countries, etc
- The structure of a phone number helps with call routing
- The phone system has evolved in ways unimagined, and phone numbers mostly work

Naming

- name in a distributed system
 - string of bits or characters that refers to an entity
- Entities
 - Hosts, printers, disks, files,
 - Processes, users, mailboxes, news groups,
 - Web pages, graphical windows, message,
 - network connections,

Names

- Identifier
 - An identifier refers to at most one entity
 - Each entity is referred by at least one identifier
 - An identifier always refers to the same entity
- Address
 - The name of an access point to an entity
- Human friendly name
 - Unix file name, DNS names
- Names are always organized in a name space
 - A name space is an organization mechanism for a group of names

Names characteristics

- Uniformity (naming/access)
- Global vs. local names
- Absolute and Relative Names
- Opacity of names (purity)
- Name resolution

Uniformity

- Structure of a name should be constant
 - In space
 - In time
- Operations should also be uniform
- Allows extensibility and openness
- Phone numbers
- DNS
- URI/URN

Global vs. local names

- Global
 - Resolve to the same value everywhere
 - Absolute name
 - Required in consistent global systems
 - Location-independent naming
- Local
 - resolution depends on context
 - Relative name
- Phone number
 - Pre-1999 – Local
 - Pos-1999 - Global

Purity / Opacity of names

- Pure /opaque names
 - The resolution algorithm does not use any information on the name
 - Don't include any clue about the address/entity
 - Allow migration of entities
 - Resolution more complex
- Impure /transparent names
 - Information on the name is used in the resolution
 - Name implies address/entity

Pure/Opaque names

- No special cases to avoid when allocating names
 - Any name usable for any object
- All names treated uniformly
 - often more efficient
- Separation of concern
 - the name solves one problem only...identifying an object

Impure / Transparent names

- It's very handy to know something about an object from its name
- Names can reflect the structure of the system:
 - /comp/150IDS/files suggests a hierarchy of files for Comp Sci courses
- If consistent patterns are used, we can often predict one name from another
 - If `http://weather.example.com/?city=miami` is the Miami weather
 - ...then maybe `http://weather.example.com/?city=boston` will get you Boston?
- Lots of different types of information winds up in names (for better or worse)
 - File type, parent directories, DNS name of host, operations
- But!
 - Only trust what the specifications say
 - In a URI, `.html` does not necessarily mean HTML, `.jpeg` need not be image/jpeg, etc.
 - `mailto:` definitely identifies a mailbox, because the specifications say it does in all cases

Purity / Opacity of names

- License plates (portugal)
 - Pure
- Phone numbers
 - Impure
- Ethernet
 - Pure
- URL
 - impure

Names / Addresses

- Name
 - an identifier for something
 - Not all names are addresses.
 - GUIDs, URNs, your name
- Address
 - a name that helps you locate something
 - HTTP URIs, “phone numbers”, postal addresses
- For names that are not addresses:
 - resolving the name to an object is necessary
 - hash table, registry, associative search, etc.

Collision / Aliasing

- Collisions
 - using one name for multiple things
- Aliasing
 - Using several names for a single thing

Aliasing

- Can be handy
- Even for simple reasons: `http://example.com` vs `HTTP://EXAMPLE.COM`
- Absolute vs. relative names for the same thing
- `a/b/c` vs `a/b/../b/c`
- Creating more than one name for an object tends to cause confusion
 - Reasoning about systems with aliases is much harder
 - Even allowing for multiple names causes trouble
 - If I have a cached copy of `http://www.example.com`, is it usable for references to `HTTP://EXAMPLE.COM`? -- cache code must be prepared to check
 - Computer languages that allow pointer aliasing are very hard to optimize...it's very hard for the compiler to keep track of which pointers might be referencing a given part of memory
 - On the Web: aliases make it hard to tell which pages are popular

Collision

- Should be allowed?
 - Usually not (no unrelated objects)
- When is it useful?
 - Multiple versions of a page that changes
 - The mobile vs. full size version of a Web page
 - A resource (a movie) vs information about the resource (blog about the movie)
- But
 - we lose the ability to reliably pick out one or the other

Should all addresses resolve

- Depends on the system
 - Some programming languages will create a variable when you access it
 - Some hypertext systems won't let you make a link unless the target exists
 - Indeed, Tim Berners-Lee's early work on the Web was rejected by academics because he didn't insist on such consistency...

Dangling links and “404” status code

- One of Tim’s key contributions was the realization that
 - enforcing such consistency is impractical on a global scale
- You can have a global Web, or you can have all links resolve...
 - but not both!
- That’s why the 404 HTTP failure code is one of the most important architectural features of the Web!

Name resolution

- The process of looking up a name
- Iterative
 - Start with the root
 - Each layer resolves as much as it can and returns address of next name server
 - Client interactively contacts servers
 - Lower performance demand on each name server
- Recursive
 - Start at the root
 - Each layer resolves as much as it can and hands the rest to the next server
 - Servers recursively contacts servers
 - Caching possible at name servers
 - Communication cost reduced

Iterative name resolution

- Start with the root
- Each layer resolves as much as it can and returns address of next name server
- Iteratively client concatenates servers

Why phone numbers?

- Good example
- Much like the web
 - Scale: worldwide
 - Time horizon: centuries
 - Evolving:
 - New devices, new operators
 - Federated
 - Distributed organization (countries, telcos)

DNS

- Primarily used for looking up host addresses and mail servers
- Comparable to a telephone book
- Its name space is hierarchically organized as a rooted tree
 - Each zone is implemented by a name server, which is always replicated
- DNS also maintains an inverse mapping of IP addresses to host names

